


20 actions → five future latent positions per view, denoised jointly

Flow tensor: 1 batch × 16 learned channels × 12 time/view slots × 32 × 32 spatial grid



Robot coverage weights only agentview slots1-5. Other slots remain unchanged.

Each spatial location contains 16 flow features; correction repeats at the selected denoising evaluations.